
Syntax and Parsing I

Constituency Parsing

Slav Petrov – Google

Thanks to:


Dan Klein, Ryan McDonald, Alexander Rush, Joakim Nivre,
Greg Durrett, David Weiss, Luheng He, Timothy Dozat

Lisbon Machine Learning School 2018

Why Parsing?

On January 13, 2018, a false ballistic missile alert was issued via the Emergency Alert System and Commercial Mobile Alert System over television, radio, and cellphones in the U.S. state of Hawaii. The alert stated that there was an incoming ballistic missile threat to Hawaii, advised residents to seek shelter, and concluded "This is not a drill". The message was sent at 8:07 a.m. local time.

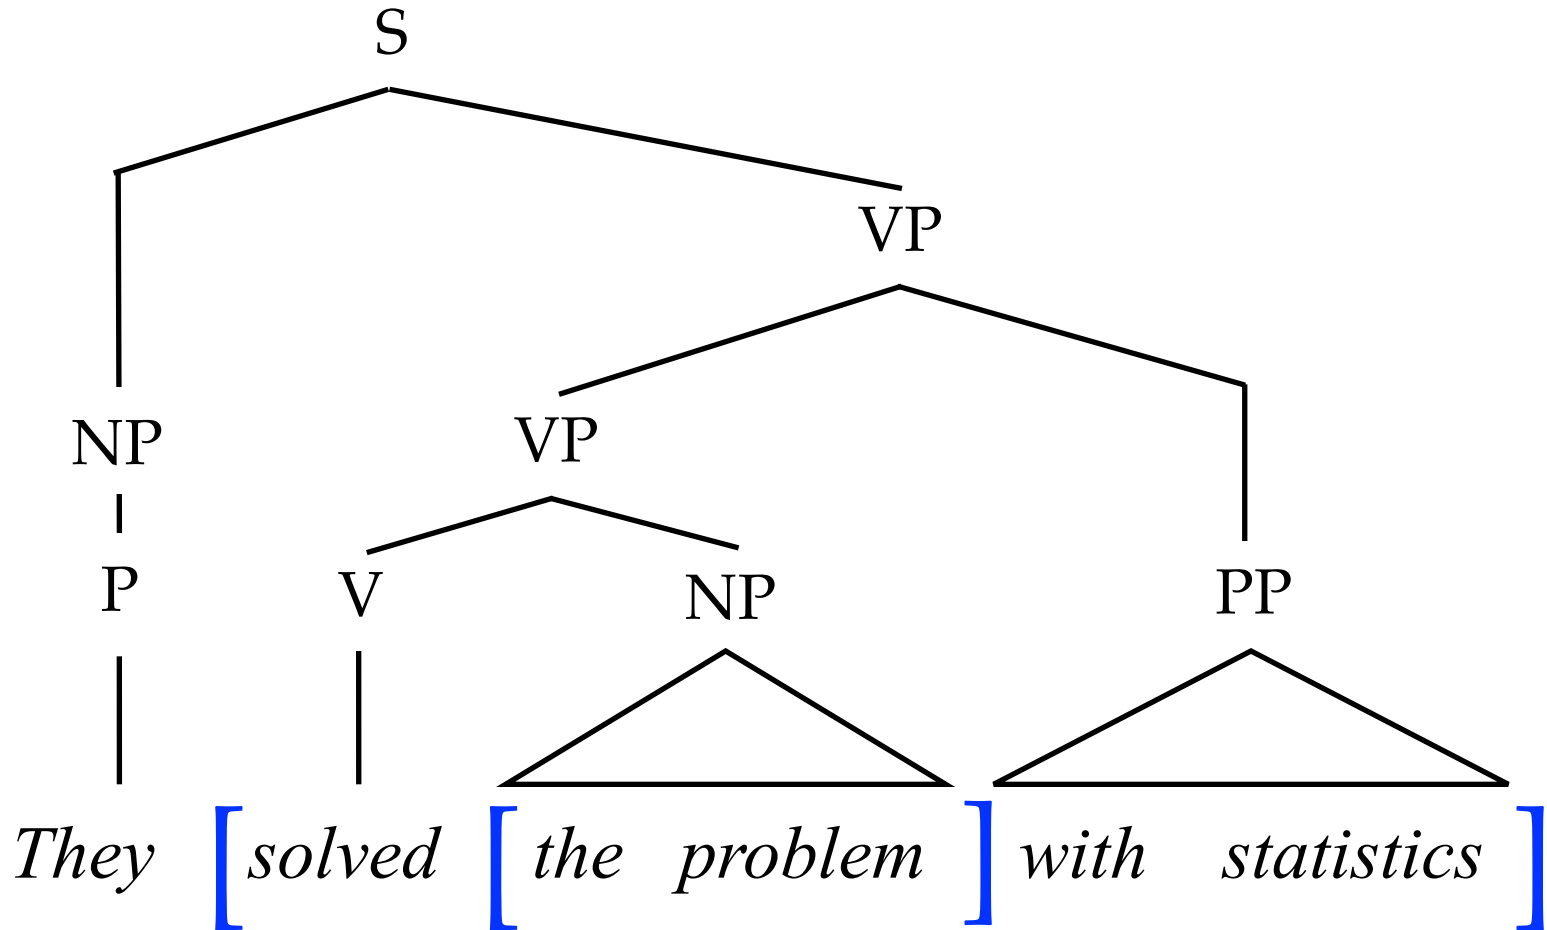
Input Document




Structure:

Syntactic,
Semantic,
etc.

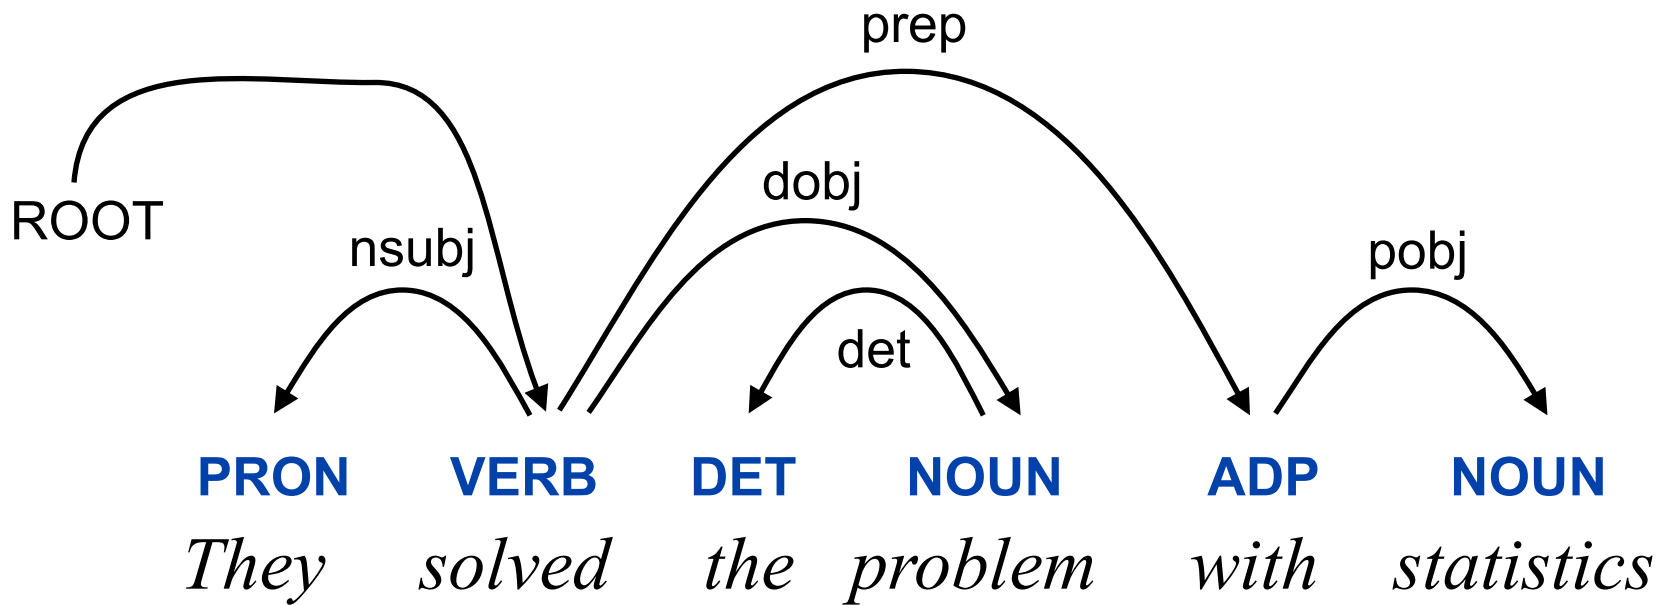
Analyzing Natural Language



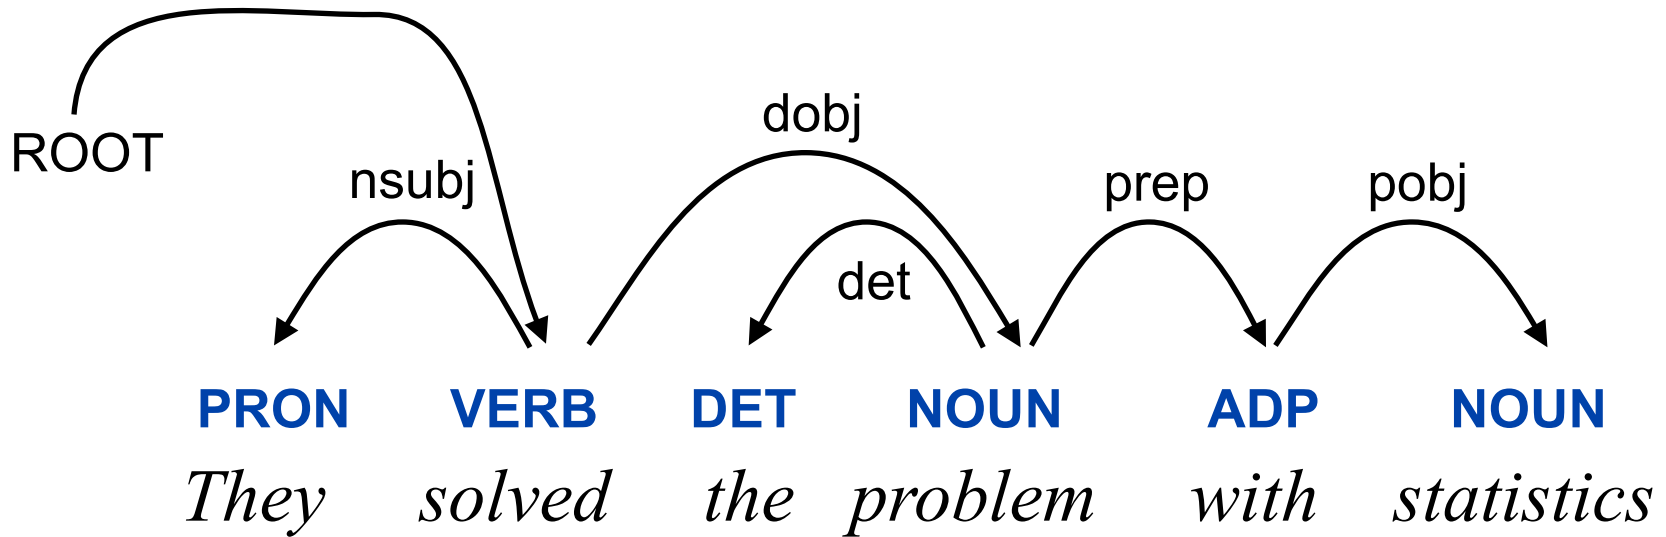
Syntax and Semantics




Constituency and Dependency



Constituency and Dependency

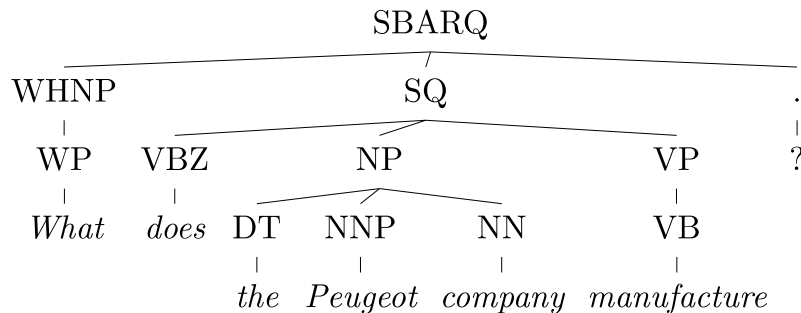


A "real" Sentence




Influential members of the House Ways and Means Committee introduced legislation that would restrict how the new savings-and-loan bailout agency can raise capital, creating another potential obstacle to the government's sale of sick thrifts.

Dependency Parsing



- Directed edges between pairs of word (head, dependent)
- Can handle free word-order languages
- Very efficient decoding algorithms exist
- Second part of today's lecture



Attachments

- I cleaned the dishes from dinner
- I cleaned the dishes with detergent
- I cleaned the dishes in my pajamas
- I cleaned the dishes in the sink

Classical NLP: Parsing

- Write symbolic or logical rules:

VBD		VB			
VBN	VBZ	VBP	VBZ		
NNP	NNS	NN	NNS	CD	NN

Fed raises interest rates 0.5 percent

Grammar (CFG)		Lexicon
ROOT → S	NP → NP PP	NN → interest
S → NP VP	VP → VBP NP	NNS → raises
NP → DT NN	VP → VBP NP PP	VBP → interest
NP → NN NNS	PP → IN NP	VBZ → raises

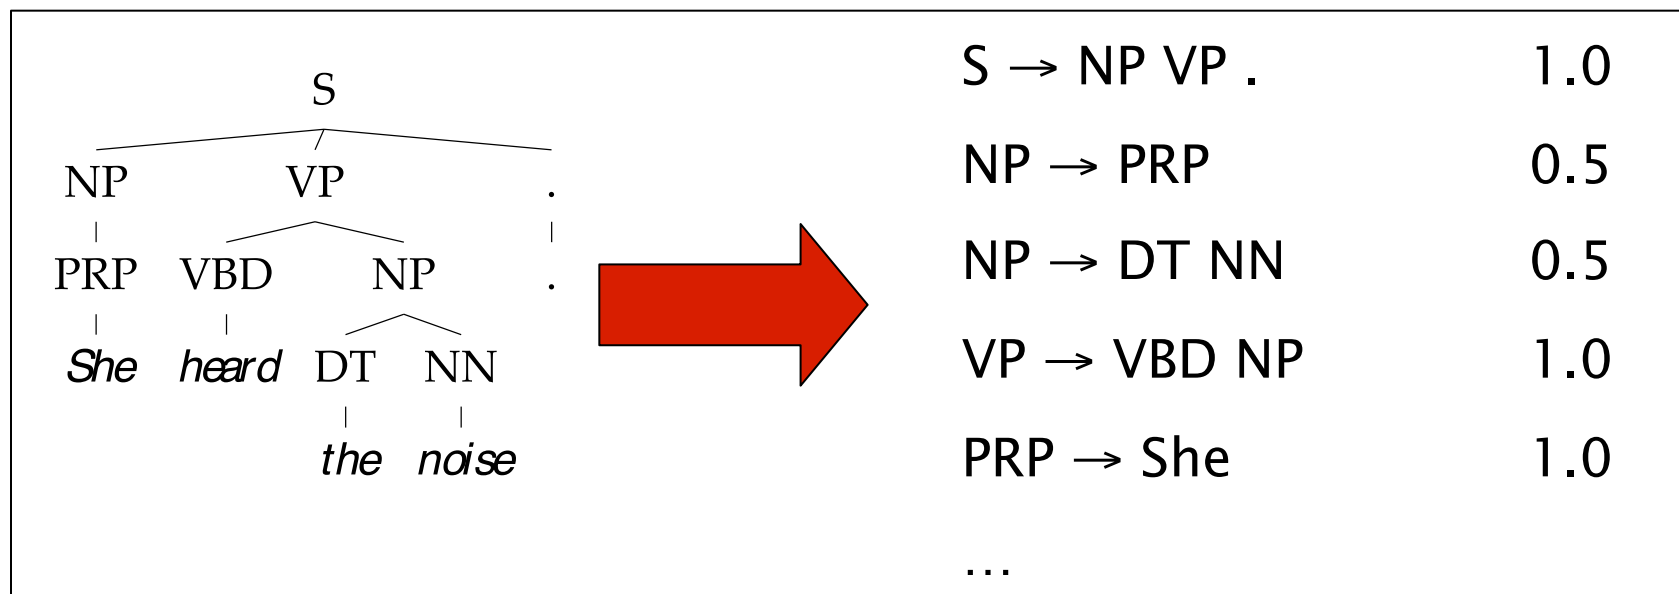
- Use deduction systems to prove parses from words
 - Minimal grammar on “Fed raises” sentence: 36 parses
 - Real-size grammar: many millions of parses
- This scaled very badly, didn’t yield broad-coverage tools

Probabilistic Context-Free Grammars

- A context-free grammar is a tuple $\langle N, T, S, R \rangle$
 - N : the set of non-terminals
 - Phrasal categories: $S, NP, VP, ADJP$, etc.
 - Parts-of-speech (pre-terminals): NN, JJ, DT, VB
 - T : the set of terminals (the words)
 - S : the start symbol
 - Often written as $ROOT$ or TOP
 - Not usually the sentence non-terminal S
 - R : the set of rules
 - Of the form $X \rightarrow Y_1 Y_2 \dots Y_k$, with $X, Y_i \in N$
 - Examples: $S \rightarrow NP VP$, $VP \rightarrow VP CC VP$
 - Also called rewrites, productions, or local trees
- A PCFG adds:
 - A top-down production probability per rule $P(Y_1 Y_2 \dots Y_k \mid X)$

Treebank Grammars

- Need a PCFG for broad coverage parsing.
- Can take a grammar right off the trees (doesn't work well):




- Better results by enriching the grammar (e.g., lexicalization).
- Can also get reasonable parsers without lexicalization.

Treebank Grammar Scale

- Treebank grammars can be enormous
 - As FSAs, the raw grammar has $\sim 10K$ states, excluding the lexicon
 - Better parsers usually make the grammars larger, not smaller


NP



Chomsky Normal Form

- Chomsky normal form:

- All rules of the form $X \rightarrow YZ$ or $X \rightarrow w$
- In principle, this is no limitation on the space of (P)CFGs
 - N-ary rules introduce new non-terminals



- Unaries / empties are “promoted”
- In practice it’s kind of a pain:
 - Reconstructing n-aries is easy
 - Reconstructing unaries is trickier
 - The straightforward transformations don’t preserve tree scores
- Makes parsing algorithms simpler!

A Recursive Parser

```
bestScore(X,i,j,s)
  if (j = i+1)
    return tagScore(X,s[i])
  else
    return max score(X->YZ) *
               bestScore(Y,i,k) *
               bestScore(Z,k,j)
```

- Will this parser work?
- Why or why not?
- Memory requirements?

A Memoized Parser

- One small change:

```
bestScore(X,i,j,s)
    if (scores[X][i][j] == null)
        if (j = i+1)
            score = tagScore(X,s[i])
        else
            score = max    score(X->YZ) *
                        bestScore(Y,i,k) *
                        bestScore(Z,k,j)
        scores[X][i][j] = score
    return scores[X][i][j]
```


A Bottom-Up Parser (CKY)

- Can also organize things bottom-up

```
bestScore(s)
```

```
  for (i : [0,n-1])
```

```
    for (X : tags[s[i]])
```

```
      score[X][i][i+1] =  
        tagScore(X,s[i])
```

```
  for (diff : [2,n])
```


```
    for (i : [0,n-diff])
```

```
      j = i + diff
```

```
      for (X->YZ : rule)
```

```
        for (k : [i+1, j-1])
```


```
          score[X][i][j] = max score[X][i][j],  
                                score(X->YZ) *  
                                score[Y][i][k] *  
                                score[Z][k][j]
```



Time: Theory

- How much time will it take to parse?

- For each diff ($\leq n$)
 - For each i ($\leq n$)
 - For each rule $X \rightarrow Y Z$
 - For each split point k
Do constant work



- Total time: $|\text{rules}| * n^3$
- Something like 5 sec for an unoptimized parse of a 20-word sentences, or 0.2sec for an optimized parser


Unary Rules

- Unary rules?

```
bestScore(X,i,j,s)
  if (j = i+1)
    return tagScore(X,s[i])
  else
    return max max score(X->YZ) *
      bestScore(Y,i,k) *
      bestScore(Z,k,j)
      max score(X->Y) *
        bestScore(Y,i,j)
```

CNF + Unary Closure

- We need unaries to be non-cyclic
 - Can address by pre-calculating the unary closure
 - Rather than having zero or more unaries, always have exactly one



- Alternate unary and binary layers
- Reconstruct unary chains afterwards

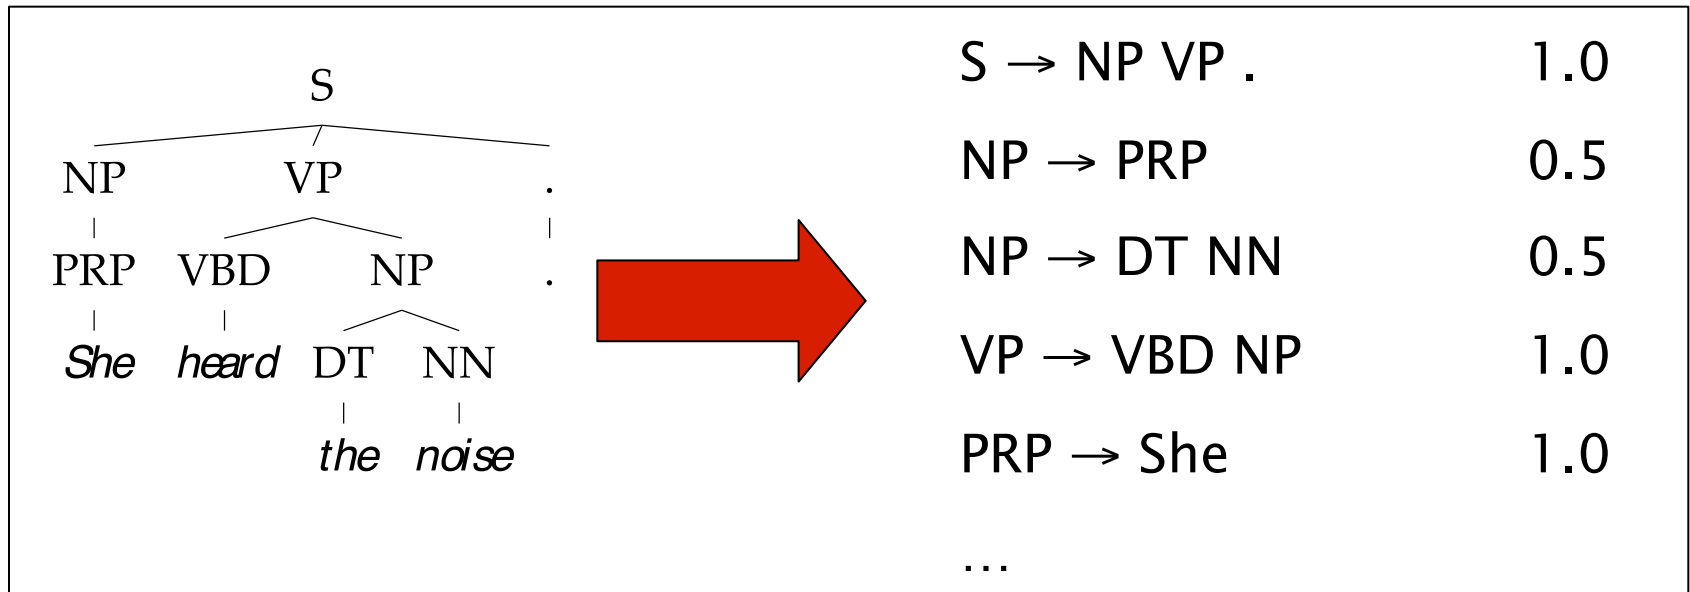
Alternating Layers

```
bestScoreB(X,i,j,s)
    return max max score(X->YZ) *
    bestScoreU(Y,i,k) *
    bestScoreU(Z,k,j)

bestScoreU(X,i,j,s)
    if (j = i+1)
        return tagScore(X,s[i])
    else
        return max max score(X->Y) *
                    bestScoreB(Y,i,j)
```

Treebank Grammars

- Need a PCFG for broad coverage parsing. [Charniak '96]
- Can take a grammar right off the trees (doesn't work well):




- Better results by enriching the grammar (e.g., lexicalization).
- Can also get reasonable parsers

<i>Model</i>	<i>F1</i>
Charniak '96	72.0

Conditional Independence?

- Not every NP expansion can fill every NP slot




- A grammar with symbols like “NP” won’t be context-free
- Statistically, conditional independence too strong


Non-Independence

- Independence assumptions are often too strong.


All NPs



NPs under S




NPs under VP



- Example: the expansion of an NP is highly dependent on the parent of the NP (i.e., subjects vs. objects).
- Also: the subject and object expansions are correlated!

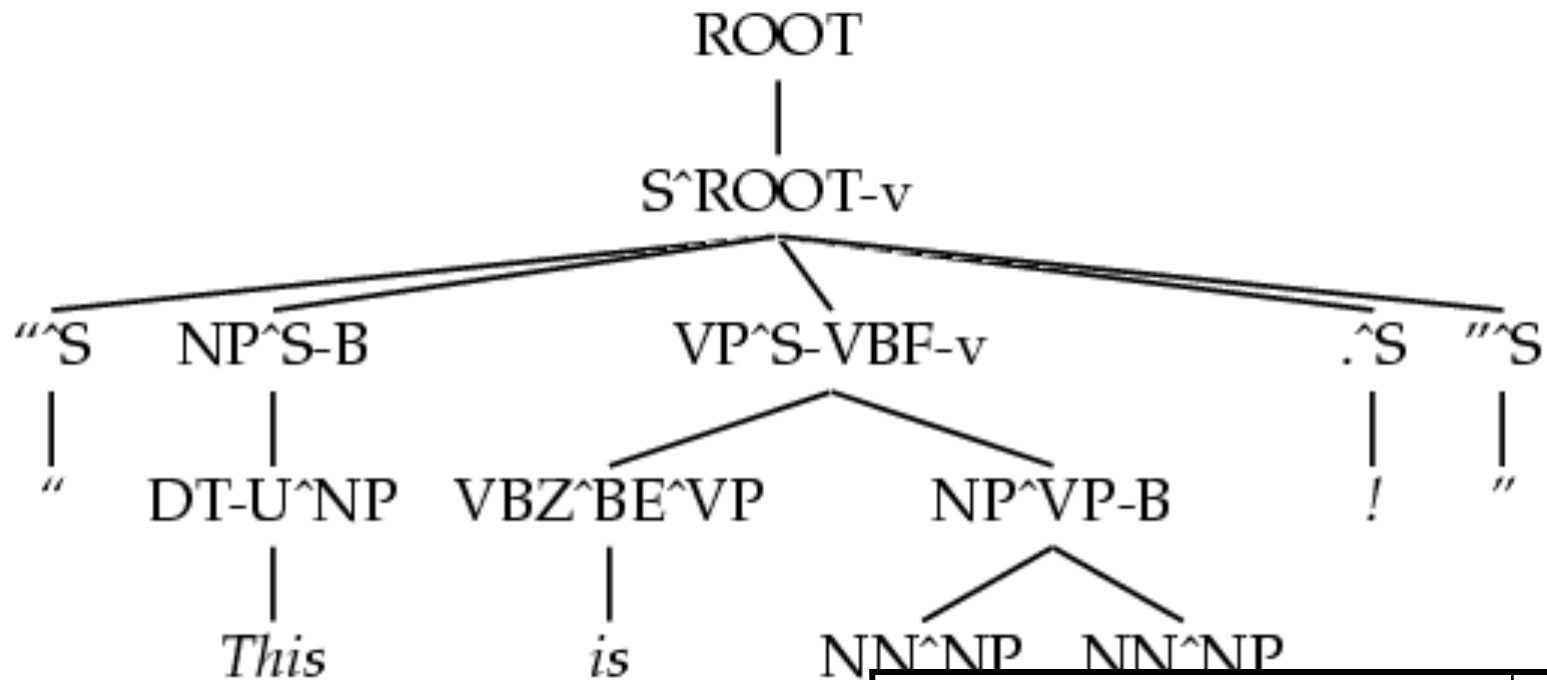
The Game of Designing a Grammar



- Structure Annotation [Johnson '98, Klein & Manning '03]
- Lexicalization [Collins '99, Charniak '00]
- Latent Variables [Matsuzaki et al. '05, Petrov et al. '06]
- (Neural) CRF Parsing [Hall et al. '14, Durrett & Klein '15]


A Fully Annotated (Unlexicalized) Tree

[Klein & Manning '03]





<i>Model</i>	<i>F1</i>
Charniak '96	72.0
Klein&Manning '03	86,3

The Game of Designing a Grammar





- Annotation refines base treebank symbols to improve statistical fit of the grammar
 - Head lexicalization [Collins '99, Charniak '00]

Problems with PCFGs



- If we do no annotation, these trees differ only in one rule:
 - $VP \rightarrow VP PP$
 - $NP \rightarrow NP PP$
- Parse will go one way or the other, regardless of words
- We addressed this in one way with unlexicalized grammars (how?)
- Lexicalization allows us to be sensitive to specific words

Problems with PCFGs



- What's different between basic PCFG scores here?
- What (lexical) correlations need to be scored?

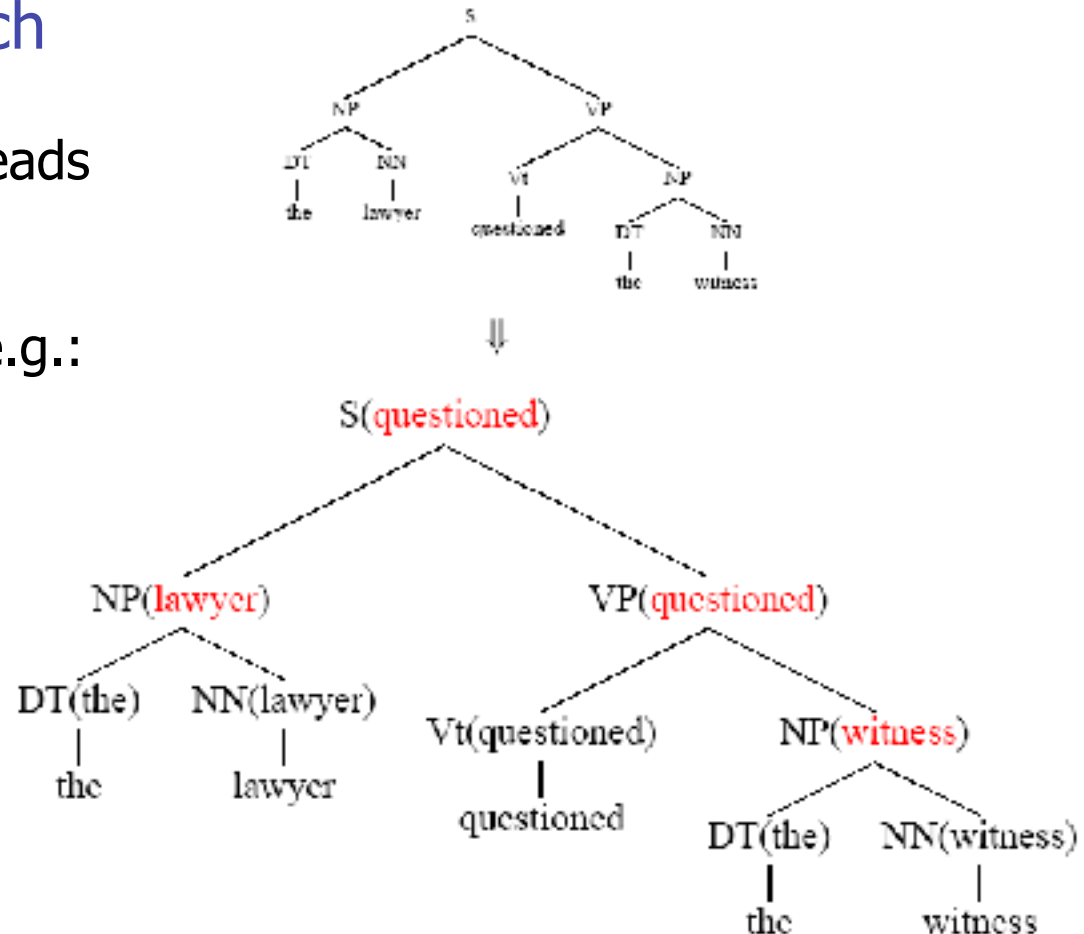
Lexicalized Trees

[Charniak '97,
Collins '97]

- Add “headwords” to each phrasal node

- Syntactic vs. semantic heads
- Headship not in (most) treebanks
- Usually use head rules, e.g.:


- NP:
 - Take leftmost NP
 - Take rightmost N*
 - Take rightmost JJ
 - Take right child
- VP:
 - Take leftmost VB*
 - Take leftmost VP
 - Take left child



Lexicalized Grammars


- Challenges:

- Many parameters to estimate: requires sophisticated smoothing techniques
- Exact inference is too slow: requires pruning heuristics
- Difficult to adapt to new languages: At least head rules need to be specified, typically more changes needed



<i>Model</i>	<i>F1</i>
Klein&Manning '03	86,3
Charniak '00	89,7


The Game of Designing a Grammar




- Annotation refines base treebank symbols to improve statistical fit of the grammar
 - Automatic clustering

Latent Variable Grammars

[Matsuzaki et al. '05, Petrov et al. '06]



Parse Tree T
Sentence w



Derivations $t : T$

Grammar G		
$S_0 \rightarrow NP_0 VP_0$?	
$S_0 \rightarrow NP_1 VP_0$?	
$S_0 \rightarrow NP_0 VP_1$?	
$S_0 \rightarrow NP_1 VP_1$?	
$S_1 \rightarrow NP_0 VP_0$?	
...		
$S_1 \rightarrow NP_1 VP_1$?	
...		
$NP_0 \rightarrow PRP_0$?	
$NP_0 \rightarrow PRP_1$?	
...		


Lexicon		
$PRP_0 \rightarrow \text{She}$?	
$PRP_1 \rightarrow \text{She}$?	
...		
$VBD_0 \rightarrow \text{was}$?	
$VBD_1 \rightarrow \text{was}$?	
$VBD_2 \rightarrow \text{was}$?	
...		

Parameters θ

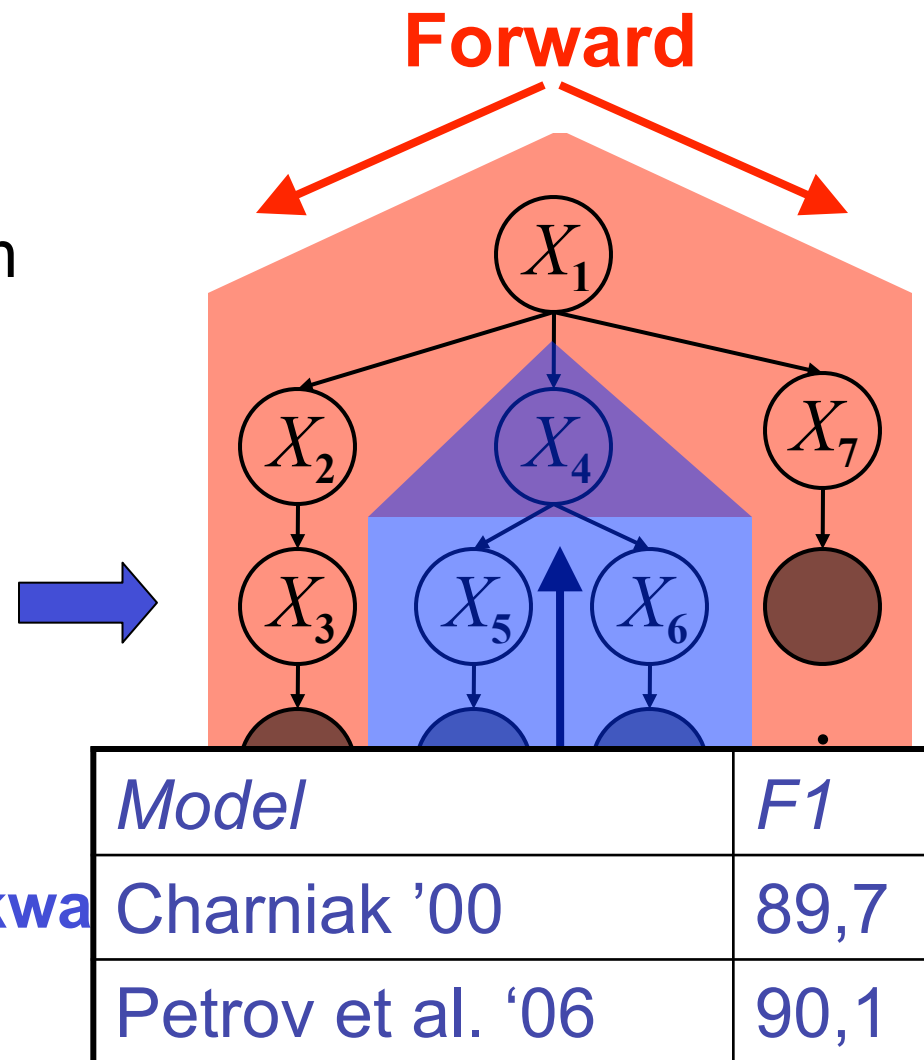
Learning Latent Annotations

EM algorithm:


- Brackets are known
- Base categories are known
- Only induce subcategories



Just like Forward-Backward
for HMMs.



The Game of Designing a Grammar




- Annotation refines base treebank symbols to improve statistical fit of the grammar
 - CRF Parsing (+Neural Network Representations)

Generative vs. Discriminative

Generative

Maximize joint likelihood
of gold tree **and** sentence




EASY: expectations over
observed trees

[Matsuzaki et al. '05, Petrov et al. '06]

Discriminative

Maximize conditional likelihood
of gold tree **given** sentence



HARD: expectations
over **all** trees

[Petrov & Klein '07, '08]

Objective Functions

Generative Objective Function:

$$\max_{\theta} \mathcal{L}_{\theta}(\text{🌳}, w_1 \dots w_n) \quad [\text{Petrov, Barrett, Thibaux \& Klein '06}]$$

Discriminative Objective Function:

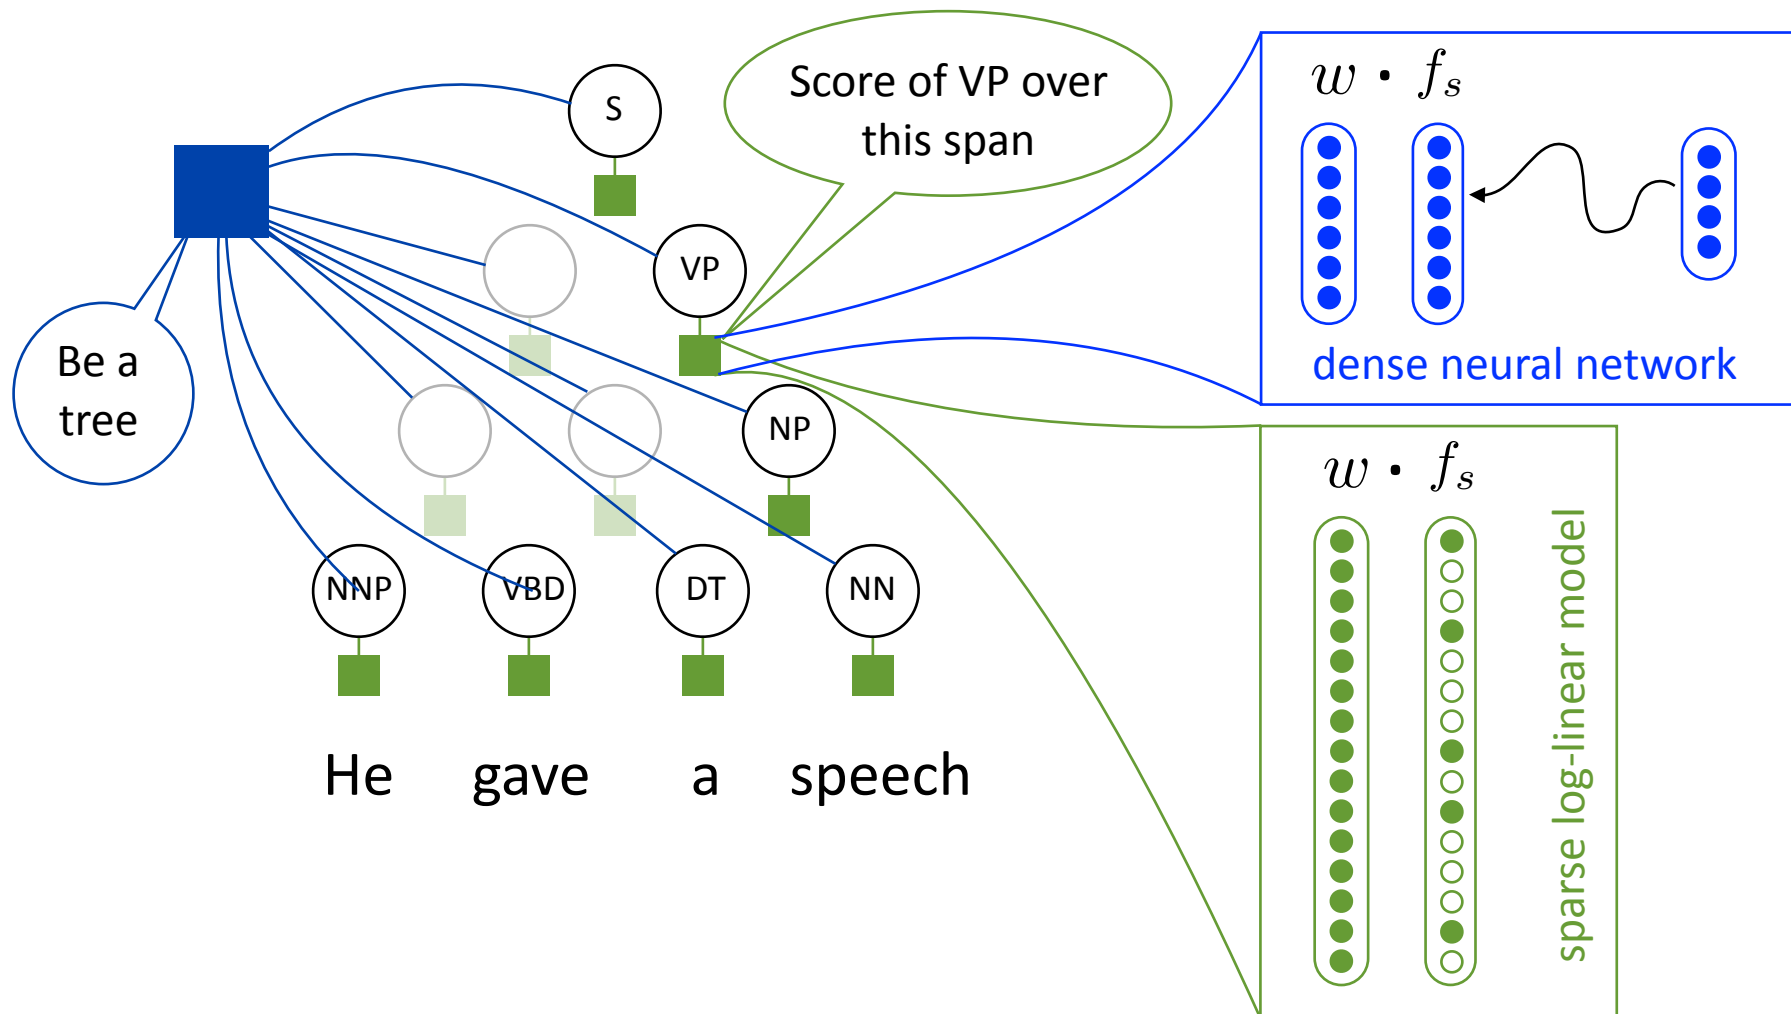
$$\max_{\theta} \mathcal{L}_{\theta}(\text{🌳} \mid w_1 \dots w_n) \quad [\text{Petrov \& Klein '08, Finkel et. al '08}]$$

Bayesian Objective Function:

$$\max_{\theta} \mathcal{P}(\theta \mid \text{🌳}) \mathcal{L}_{\theta}(\text{🌳}, w_1 \dots w_n) \quad [\text{Liang, Petrov, Jordan \& Klein '07}]$$

(Neural) CRF Parsing

[Taskar et al. '04. Petrov & Klein '07, Hall et al. '14, Durrett et al. '15]



CRF Parsing Sparse Features

$$P(T|x) \propto \prod_{r \in T} \exp(\text{score}(r))$$

$$\text{score}({}_2\text{NP}_7 \rightarrow {}_2\text{NP}_4 {}_4\text{PP}_7) = w^\top f({}_2\text{NP}_7 \rightarrow {}_2\text{NP}_4 {}_4\text{PP}_7)$$

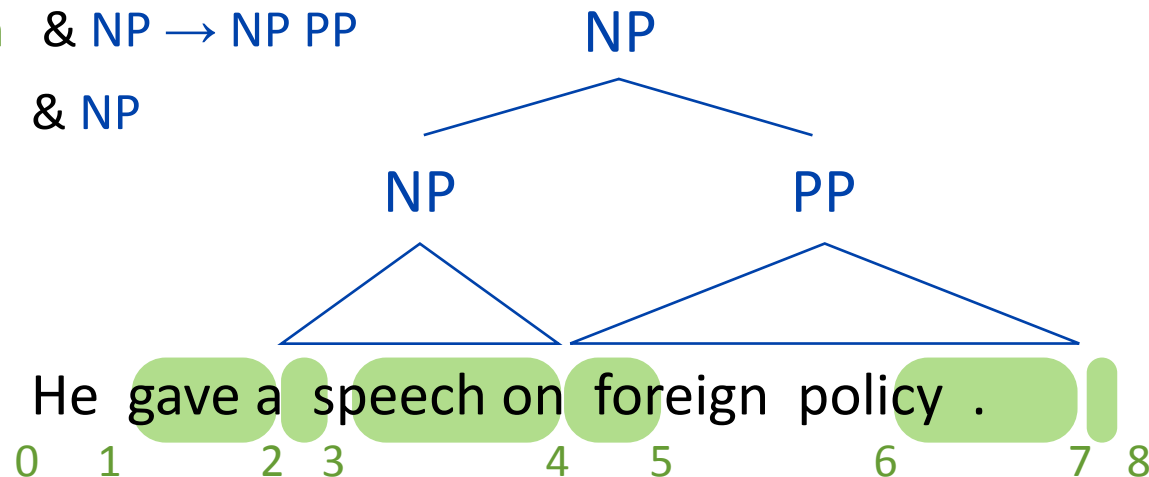
FirstWord = a & NP → NP PP

PrevWord = gave & NP → NP PP

AfterSplit = on & NP → NP PP

FirstWord = a & NP

...




Neural CRF Model

[Durrett et al. '15]

score($_2\text{NP}_7 \rightarrow _2\text{NP}_4 \text{ } _4\text{PP}_7$) =

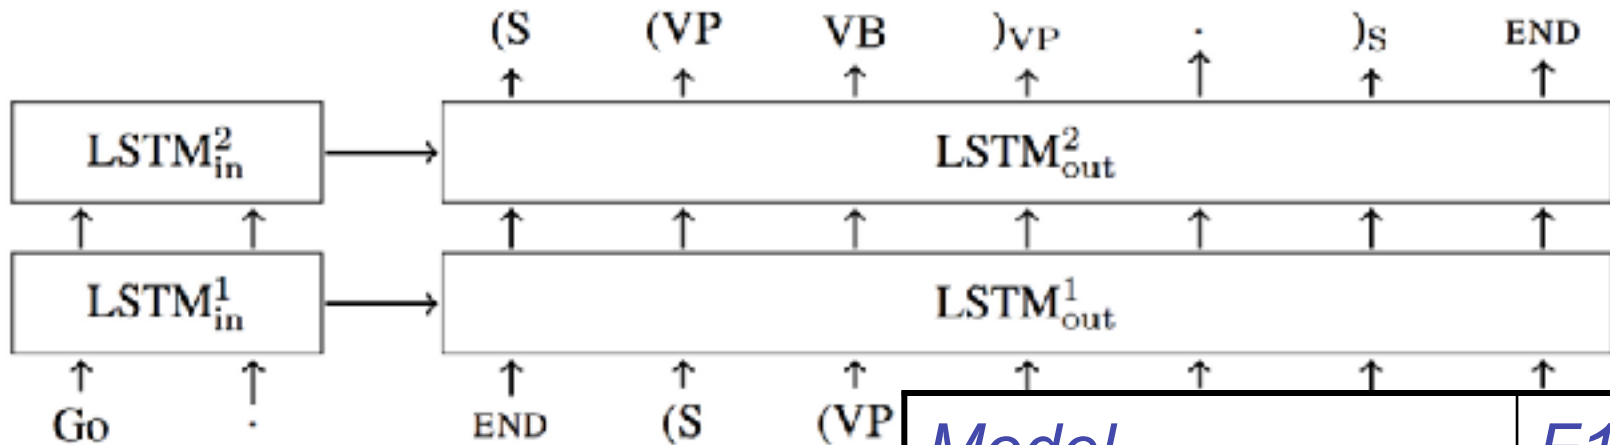
$$W \odot \left(f_s \left(_2\text{X}_7 \rightarrow _2\text{X}_4 \text{ } _4\text{X}_7 \right) f_o^\top (\text{NP} \rightarrow \text{NP PP}) \right)$$



<i>Model</i>	<i>F1</i>
Petrov et al. '06	90,1
Durrett et. al. '15	91,1

LSTM Parsing [Vinyals et al. '15]

- Treat parsing as a sequence-to-sequence prediction problem
- Completely ignores tree structure, uses LSTMs as black boxes
- No global normalization, only local normalization




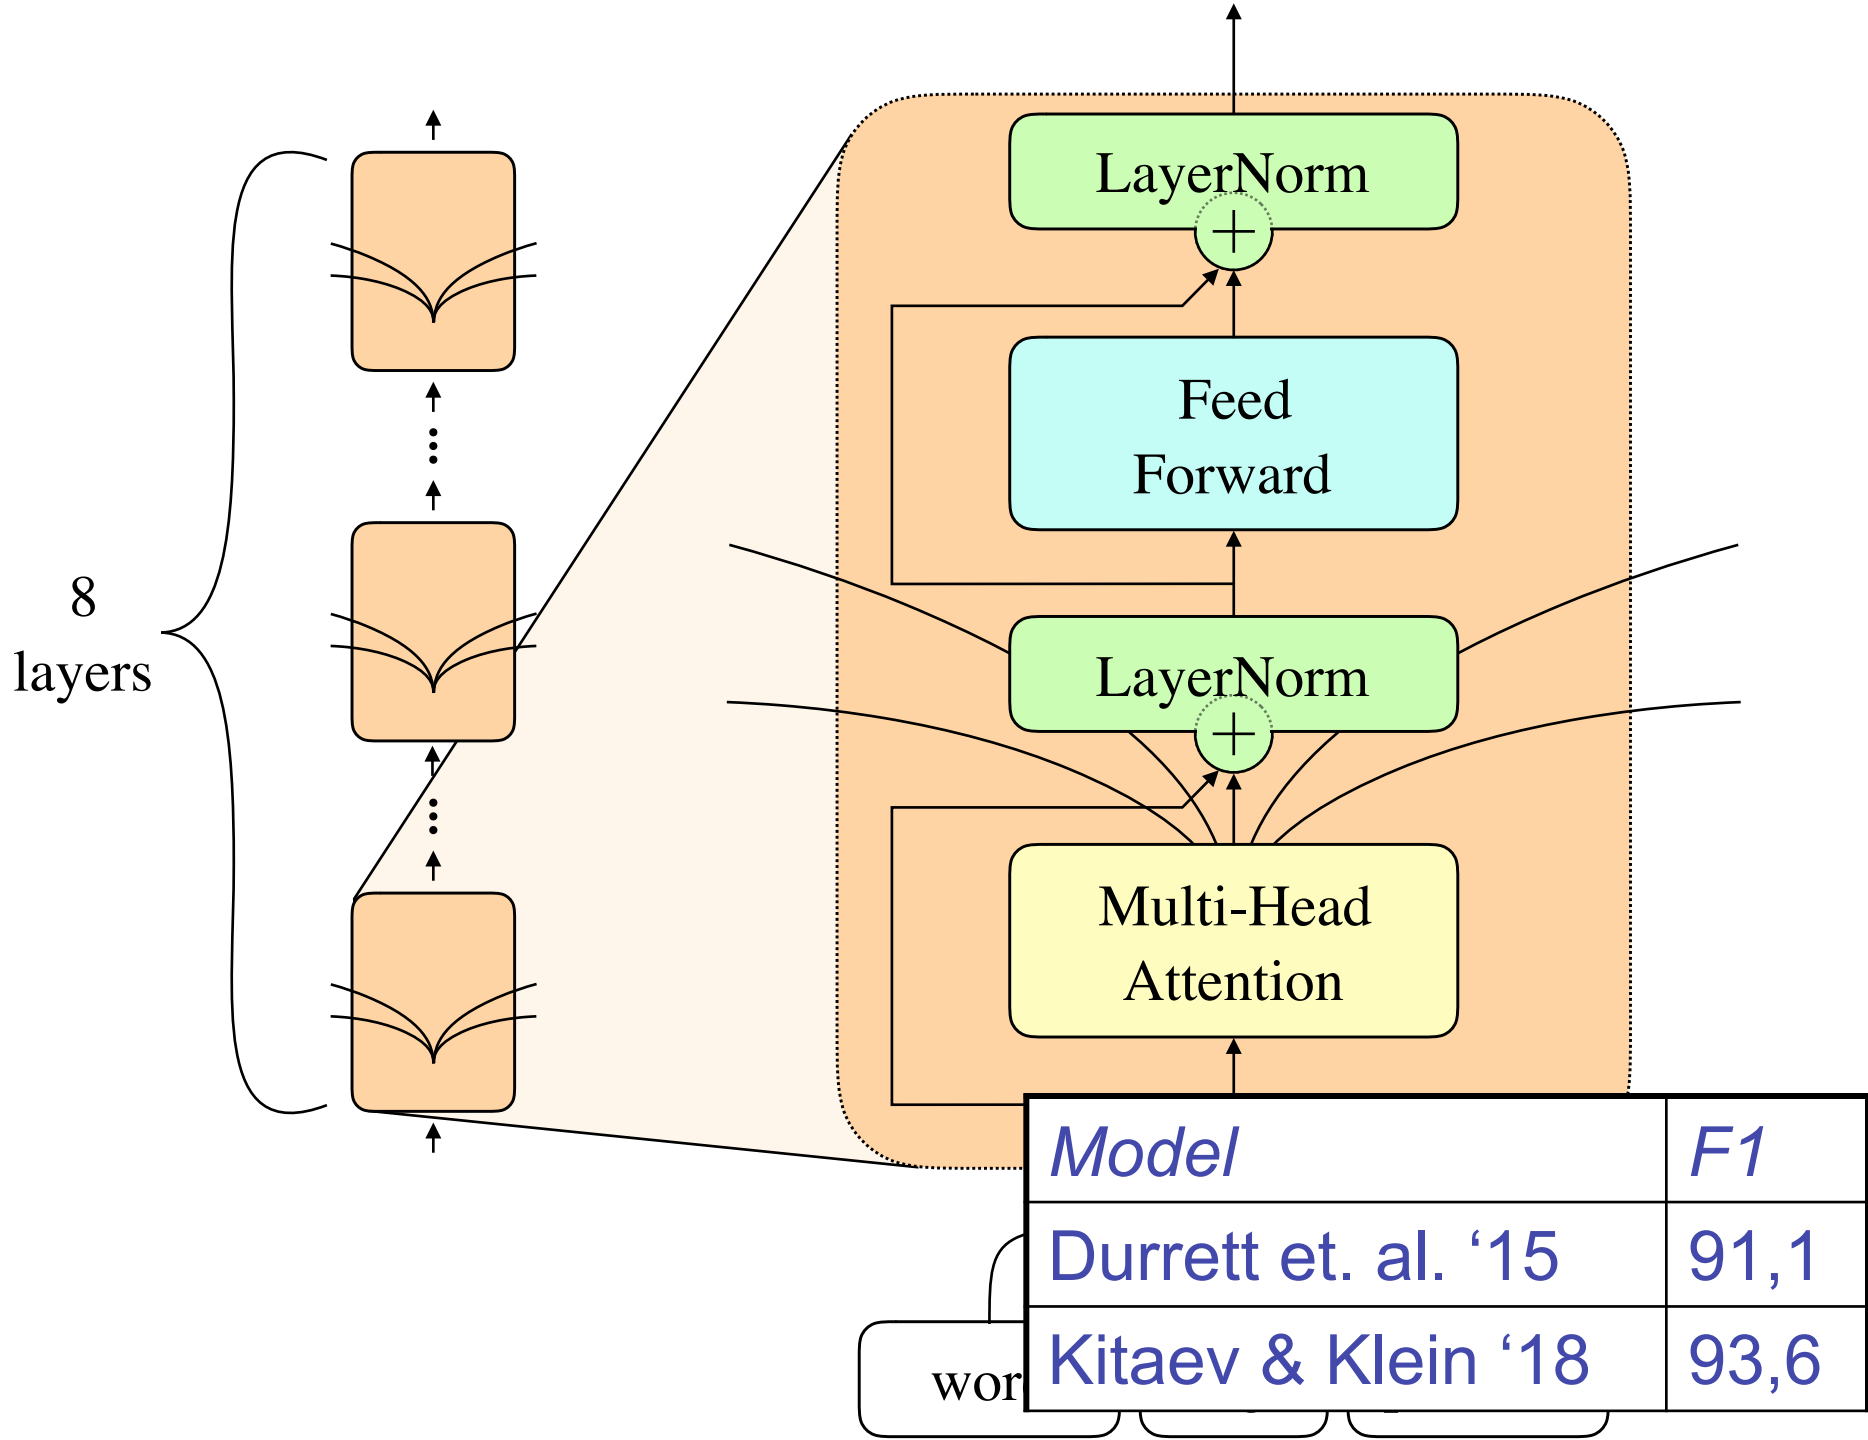
$$P(y_1, \dots, y_{T'} | x_1, \dots, x_T) = \prod_{t=1}^{T'} P(y_t | x_1, \dots, x_T, y_1, \dots, y_{t-1})$$

<i>Model</i>	<i>F1</i>
Durrett et. al. '15	91,1
Vinyals et al. '15	88.6*


Parsing with Self-Attention

[Kitaev & Kleint '18]






Detailed English Results (Old)



Detailed English Results (New)



Multi-Lingual Results

